

OMNIEDUBENCH: A COMPREHENSIVE CHINESE BENCHMARK FOR EVALUATING LARGE LANGUAGE MODELS IN EDUCATION

Min Zhang^{1*} **Hao Chen¹** **Hao Chen¹** **Wenqi Zhang²** **Didi Zhu³** **Xin Lin¹**
Bo Jiang^{1†} **Aimin Zhou^{1†}** **Fei Wu²** **Kun Kuang²**

¹East China Normal University ²Zhejiang University ³Imperial College London

mzhang@cs.ecnu.edu.cn bjiang@deit.ecnu.edu.cn amzhou@cs.ecnu.edu.cn

<https://mind-lab-ecnu.github.io/OmniEduBench/>

ABSTRACT

With the rapid development of large language models (LLMs), various LLM-based works have been widely applied in educational fields. However, most existing LLMs and their benchmarks focus primarily on *the knowledge dimension*, largely neglecting the evaluation of cultivation capabilities that are essential for real-world educational scenarios. Additionally, current benchmarks are often limited to a single subject or question type, lacking sufficient diversity. This issue is particularly prominent within the Chinese context. To address this gap, we introduce **OmniEduBench, a comprehensive Chinese educational benchmark**. OmniEduBench consists of 24.602K high-quality question-answer pairs. The data is meticulously divided into two core dimensions: **the knowledge dimension and the cultivation dimension**, which contain 18.121K and 6.481K entries, respectively. Each dimension is further subdivided into 6 fine-grained categories, covering a total of 61 different subjects (41 in the knowledge and 20 in the cultivation). Furthermore, the dataset features a rich variety of question formats, including 11 common exam question types, providing a solid foundation for comprehensively evaluating LLMs' capabilities in education. Extensive experiments on 11 mainstream open-source and closed-source LLMs reveal a clear performance gap. In the knowledge dimension, only Gemini-2.5 Pro surpassed 60% accuracy, while in the cultivation dimension, the best-performing model, QWQ, still trailed human intelligence by nearly 30%. These results highlight the substantial room for improvement and underscore the challenges of applying LLMs in education.

1 INTRODUCTION

With the rapid emergence of large language models (LLMs), evaluation benchmarks have become increasingly critical, shifting the focus of assessment toward broader and complex skills. To address the demands of this complex paradigm, a variety of benchmarks have been proposed to evaluate the diverse capabilities of LLMs. These benchmarks cover a wide spectrum of areas, including knowledge and language understanding (*e.g.*, MMLU (Hendrycks et al., 2021), ARC (Clark et al., 2018)), reasoning (*e.g.*, GSM8K (Cobbe et al., 2021), AIME (Patel et al., 2024)), multi-turn open-ended dialogue (*e.g.*, MT-bench (Bai et al., 2024)), and coding (*e.g.*, MBPP (Austin et al., 2021)). Serving as indispensable tools for advancing LLM development, these benchmarks have been widely adopted in recent influential works (Hurst et al., 2024; Liu et al., 2024a; Team, 2025b; Comanici et al., 2025; Taylor et al., 2022; Touvron et al., 2023; OpenAI, 2023; Hoffmann et al., 2022).

In recent years, a series of powerful Chinese LLMs emerged, such as the Qwen (Yang et al., 2024; 2025), DeepSeek (Liu et al., 2024a; Guo et al., 2025; Liu et al., 2024b), achieving performance levels comparable to overseas LLMs. With the growing application of LLMs in education, researchers

*Project leader. Min Zhang (mzhang@cs.ecnu.edu.cn)

†Corresponding authors. Bo Jiang (bjiang@deit.ecnu.edu.cn) and Aimin Zhou (amzhou@cs.ecnu.edu.cn)

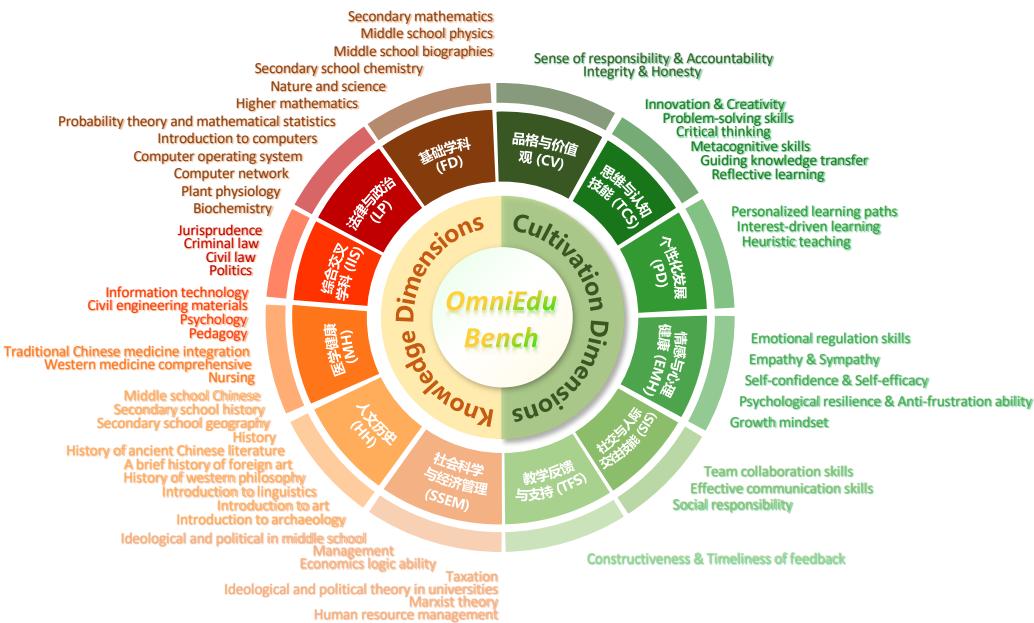


Figure 1: Overview of OmniEduBench. The benchmark comprises two dimensions: 41 subjects across six categories in the knowledge, and 20 subjects across six categories in the cultivation.

have also begun to propose Chinese education benchmarks, which can be broadly categorized into two types: (1) datasets translated from other languages and (2) datasets natively constructed from Chinese education corpora. Specifically, (1) Datasets translated from other languages refer to benchmarks constructed by directly translating existing benchmarks from other languages into Chinese. A representative work is CLUE (Xu et al., 2020a), which was translated from the English GLUE (Wang et al., 2019). However, a simple translation approach is insufficient for a rigorous evaluation of LLMs in Chinese. These datasets often fail to reflect the unique linguistic and cultural challenges of the Chinese education and inherently carry biases from their original environment, thus limiting their ability to assess LLMs' understanding of local education knowledge and teacher-student needs.

(2) Datasets natively constructed from Chinese educational corpora refer to benchmarks directly collected from Chinese educational text resources, such as C-Eval (Huang et al., 2023), Edubench (Xu et al., 2025), Scieval (Sun et al., 2024), AGIEval (Zhong et al., 2023), and SuperCLUE (Xu et al., 2023). However, most existing education benchmarks are often limited to a single subject or question type, lacking sufficient diversity. Additionally, these datasets typically focus on the knowledge dimension, overlooking the unique cultivation aspects that are essential in real-world education.

We present OmniEduBench, a comprehensive Chinese education benchmark designed to thoroughly evaluate LLMs in terms of both knowledge understanding and skill cultivation in educational scenarios. OmniEduBench encompasses knowledge and cultivation dimension and comprises a total of 24.602K high-quality question-answer pairs, covering 11 common exam question types (e.g., multiple choice (单选题), multiple answer (多选题), fill-in-the-blank (填空题), short answer (简答题), composite questions (复合题), term explanation (名词解释), True/False (判断题), calculation (计算题), logical reasoning (逻辑推理题), case analysis (案例分析题), and essay (论述题)), as illustrated in Figure 1. The knowledge dimension includes 18.121K question-answer pairs spanning 41 subject areas, from humanities to science and engineering, and covering five difficulty levels: elementary school, middle school, high school, college, and professional examinations. The cultivation dimension comprises 6.481K question-answer pairs across 20 teaching-related comments, including guided teaching, student emotional support, and moral education (see Table 1 for details), aiming to comprehensively assess the diverse competencies required in real-world educational settings. Extensive experiments demonstrate that our proposed OmniEduBench presents a highly challenging and significant benchmark for Chinese educational evaluation. Additionally, we introduce OmniEduBench HARD, a high-difficulty subset of OmniEduBench, specifically targeting particularly demanding subjects such as advanced mathematics and competitions that require sophisticated reasoning skills. Even the state-of-the-art LLMs achieve less than 50% accuracy on this subset, highlighting the rigor and necessity of our proposed OmniEduBench education benchmark.

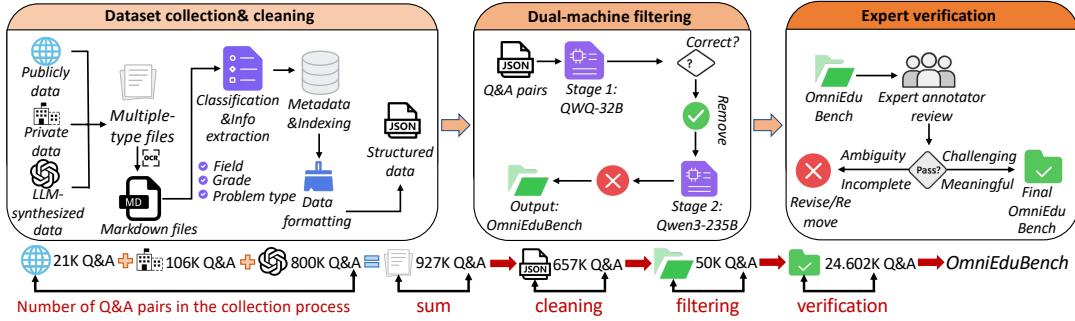


Figure 2: Overview of the construction process, including collection, cleaning, filtering, verification.

2 OMNIEDUBENCH

Our proposed OmniEduBench education benchmark is designed as a **natively Chinese education evaluation benchmark** that captures the unique linguistic and cultural knowledge of Chinese education, encompasses diverse question types, and assesses LLMs not only on their knowledge capabilities but also on the distinctive cultivation competencies required in real-world educational scenarios.

2.1 TASK DEFINITION

Knowledge dimension focuses on evaluating the model’s mastery of subject-specific knowledge. Tasks in this dimension include 11 common exam question types (*e.g.*, multiple choice, multiple answer, fill-in-the-blank, short answer, composite questions, term explanation, True/False, calculation, logical reasoning, case analysis, and essay). These 11 question types span a wide range of disciplines, from humanities and history to science, engineering, and professional fields. The primary goal is to assess the LLM’s problem-solving capabilities within the context of real-world education.

Cultivation dimension assesses LLMs on their ability to support holistic educational objectives beyond mere knowledge acquisition. This includes guiding students’ thinking processes, fostering moral and value development, enhancing emotional understanding, and promoting critical reasoning skills. Tasks in this dimension are designed to reflect realistic learning scenarios, where models must provide pedagogically sound feedback that aligns with students’ cognitive and emotional needs.

2.2 BENCHMARK CONSTRUCTION

In this section, we provide a detailed overview of the construction process for the proposed OmniEduBench education evaluation benchmark, as illustrated in Figure 2. The process consists of four key stages: dataset collection, dataset cleaning, dual-machine filtering, and expert verification.

Dataset collection. OmniEduBench is designed to encompass a wide range of diverse scenarios to enable comprehensive evaluation. To achieve this, we employ three distinct data collection methods, carefully balancing diversity and efficiency in the construction of the OmniEduBench benchmark.

Manual collection of publicly available data. Existing benchmarks often lack sufficient diversity in question types and knowledge coverage, making them inadequate for our 41 subjects in knowledge dimensions. To address this gap, we manually collected additional data from publicly available online resources (*e.g.*, XuekeNet, ZujuanNet, ShijuanNet, ShitiNet) to enrich diversity and ensure coverage of underrepresented scenarios, such as primary and career education. Furthermore, guided by the Catalogue of Undergraduate Programs in Regular Higher Education Institutions¹ issued by China’s Ministry of Education, we curated a large body of review materials and exam questions across 13 academic disciplines, including philosophy, education, law, literature, history, science, engineering, agriculture, medicine, military science, management, and the arts. This effort significantly improves distributional balance and provides a more faithful reflection of real applications.

¹http://www.moe.gov.cn/srcsite/A08/moe_1034/s4930/202403/W020240319305498791768.pdf

Manual collection of private data. Data contamination remains one of the most critical challenges in constructing evaluation datasets for LLMs. To mitigate this risk, we manually collected additional data from private resources, such as internal school exam papers. Unlike widely circulated national exams, these materials have never appeared on the public Internet or been included in large-scale web crawls, effectively reducing the risk of leakage. Incorporating such private data enhances the reliability and fairness of the benchmark, while providing a more rigorous assessment of models.

LLM-generated data. Given the difficulty of directly obtaining data in the cultivation dimension, we leveraged LLMs to generate a substantial number of scenario-based question–answer pairs, aiming to supplement gaps in existing resources. To ensure the quality of the synthetic data, we invited five education experts to conduct discussions on 20 cultivation subjects and consulted relevant books, papers, and other materials. The collected content was organized into a database, which was then provided to the LLM to enhance the fidelity and accuracy of the generated data. For generated questions, to increase their challenge, we designed highly confounding distractors via prompts and conducted sampling checks and revisions with expert verification (please see more details in expert verification). Finally, we collected a total of 927K question–answer (Q&A) pairs, including 21K from publicly available data, 106K from private data, and 800K generated by LLMs.

Dataset cleaning. The entire data cleaning process consists of multiple steps. (1) We used MinerU (Wang et al., 2024) to convert the collected 927K Q&A pairs into Markdown (md) format, enabling structured management and efficient information extraction. (2) Detailed metadata were extracted for each question, including subject, grade level, question type, and knowledge tags, to construct comprehensive question profiles that facilitate data management and subsequent analysis. (3) Standard data cleaning procedures were applied, including deduplication, removal of questions with missing key content, filtering of sensitive or inappropriate content, and exclusion of questions that rely on external information. After the cleaning process, we obtained a total of 657K Q&A pairs.

Dual-machine filtering. To ensure OmniEduBench is a high-quality and challenging benchmark, we implemented a dual-model filtering mechanism on an initial set of 657K Q&A pairs. Specifically, we first evaluated all questions using QWQ32B (Team, 2025c), retaining only those that the model answered incorrectly. This initial filtering resulted in a subset of 430K Q&A pairs. These questions then underwent a second filtering stage with the same strategy, this time using Qwen3-235B (Yang et al., 2025), ultimately yielding the final set of 50K high-quality and challenging data.

Expert verification. We recruited 50 master’s students to perform an initial quality check on the dataset based on five predefined dimensions (as shown in Table 2), removing any data that did not meet the criteria, which resulted in a final set of 24.602K Q&A pairs for OmniEduBench. Subsequently, we invited 5 senior annotation experts to conduct a rigorous quality review on a 15% random sample of the OmniEduBench. The review results, shown in Table 2, indicate that the dataset maintains high overall quality, demonstrating both reliability and applicability.

2.3 EVALUATION CRITERIA

Based on the characteristics of different question types, we adopt two evaluation metrics: (1) **Choice**. For questions with a standard answer, we directly evaluate the provided answer. This simplifies the scoring process, as the model only needs to select the most appropriate option, thereby reducing ambiguity in assessment. (2) **LLM-assisted scoring**. For short-answer questions that may have multiple valid forms but are semantically equivalent, we employ an LLM-assisted scoring method. This approach provides greater flexibility, avoids imposing unnecessary constraints on the model, and allows for a more accurate evaluation of the model’s semantic understanding and expression.

2.4 STATISTICS

Through rigorous data filtering and expert validation, we collected 18.121K high-quality question–answer pairs for the knowledge and 6.481K for the cultivation. As illustrated in Figure 1 and summarized in Table 1, with more detailed per-subject statistics provided in the Appendix, the dataset spans 12 major categories, as shown in Table 1, including K-12, higher school, university-level courses, and cultivation aspects such as emotion and reasoning, covering a total of 61 specific scenarios. Figures 3 and 4 present some representative examples in different dimensions and ques-

Table 1: Statistics of OmniEduBench and more detailed per-subject information are shown in the Appendix. Bilingual names and abbreviations of six knowledge and six cultivation dimensions.

Knowledge dimension			Cultivation dimension					
English name	Abbreviation	Chinese name	English name	Abbreviation	Chinese name			
Law & Politics	LP	法律与政治	Character & Values	CV	品格与价值观			
Foundational Disciplines	FD	基础学科	Personalized Development	PD	个性化发展			
Humanities & History	HH	人文与历史	Social & Interpersonal Skills	SIS	社会与人际交往			
Medicine & Health	MH	医学与健康	Thinking & Cognitive Skills	TCS	思维与认知能力			
Interdisciplinary & Integrated Subjects	IIS	综合与交叉学科	Teaching Feedback & Support	TFS	教学反馈与支持			
Social Sciences & Economics Management	SSEM	社会科学与经济管理	Emotional & Mental Health	EMH	情感与心理健康			
Category	Subjects	Questions	Category	Subjects	Questions			
<i>In terms of dimension</i>		<i>In terms of Knowledge</i>		<i>In terms of Cultivation</i>				
Knowledge	41	18,121	LP	4	1,455	CV	2	694
Cultivation	20	6,481	FD	11	7,918	PD	3	1,031
<i>In terms of different level</i>		HH		5,331		SIS		736
K-12 Schools	10	4,384	MH	3	918	TCS	6	1,900
High school	11	6,735	IIS	4	914	TFS	1	193
College	30	6,364	SSEM	9	1,643	EMH	5	1,833
Total	61	24,602	Total	41	18,179	Total	20	6,387

Table 2: Expert validation results for the OmniEduBench dataset.

Metric English name	Metric Chinese name	Average	Standard deviation	Inter-rater agreement
Overall quality	整体质量	4.8	0.1	0.90
Clarity	问题清晰度	4.5	0.2	0.85
Option perplexity	选项困惑度	4.8	0.3	0.83
Accuracy	答案准确性	4.8	0.1	0.90
Cultivation value	育人价值	4.6	0.2	0.88

tion types. The questions exhibit wide variability in type and difficulty and are sourced from diverse origins, primarily newly collected from public or private resources or manually constructed.

3 EXPERIMENTS

In this section, we evaluate the performance of state-of-the-art (SOTA) methods in both English and Chinese. The experimental results indicate that OmniEduBench remains a competitive benchmark.

3.1 EXPERIMENTAL SETUP

Baselines. We evaluate 11 mainstream large language models (LLMs) in total, including 3 cutting-edge closed-source models and 8 open-source models, one of which is a newly released education-oriented model. The closed-source models are GPT-4o (Hurst et al., 2024), Gemini-2.5 Pro (Comanici et al., 2025), and Claude-4 Sonnet (Team, 2025a). For the open-source models, we consider two main factors. First, they are grouped by parameter size into small (8B), medium (14B/32B/36B), and large (72B/235B/671B) scales. Second, they are categorized by functionality into: (a) general

下列说法错误的是：
The following statements are incorrect:

- A. 甲苯与甘油燃烧时产生的水量在等物质的量条件下保持恒定
A. The amount of water produced by the combustion of toluene and glycerol remains constant under the condition of equal amounts
 - B. 溴乙烷发生消去反应后，可用酸性高锰酸钾溶液直接检验生成的乙烯
B. After the elimination reaction of ethyl bromide, the ethylene produced can be directly tested with acidic potassium permanganate solution
 - C. 酰类物质在碱性条件下与银氨溶液共热可发生银镜反应
C. Aldehydes can undergo silver mirror reaction when co-heated with silver ammonia solution under alkaline conditions
 - D. 碳酸钠溶液可溶解苯酚生成苯酚钠和碳酸氢钠
D. Sodium carbonate solution dissolves phenol to produce sodium phenol and sodium bicarbonate
- 答案: B
Answer: B

(a) Single-choice question type within the knowledge dimension

有学生在参观烈士陵园时嬉笑打闹，我很生气，该怎么处理？
Some students laughed and played while visiting the martyrs' cemetery, I feel very angry, how should I deal with it?

- A. 当场严厉批评，让他道歉并写检讨。
A. Criticize harshly on the spot, ask him to apologize and write a review.
 - B. 暂停他的课外活动资格一周作为惩罚。
B. Suspension of his extracurricular activities for one week as punishment.
 - C. 课后找他谈话：你知道我们为什么要来这里吗？你觉得这样的行为对得起先烈吗？你希望别人怎么纪念你？
C. Talk to him after class: Do you know why we are here? Do you think such behavior is worthy of the martyrs? How do you want others to remember you?
 - D. 让他下次参观时担任讲解员，负责介绍一位烈士的事迹。
D. Let him serve as a docent on his next visit, responsible for introducing the deeds of a martyr.
- 答案: D
Answer: D

(b) Single-choice question type within the cultivation dimension

Figure 3: Example of (a) a single-choice question in the knowledge from a college chemist. (b) A single-choice question in the cultivation. English translations are shown for better readability.

BAT (棕色脂肪组织) 细胞膜上具有下列哪些受体? (多选)
 Which of the following receptors are present on the BAT (brown adipose tissue) cell membrane? (Multiple Choice)

A. 葡萄糖受体
 A. Glucose receptors
 B. 神经递质受体
 B. Neurotransmitter receptors
 C. 促胰液素受体
 C. Secretin-stimulating receptors
 D. 甲状腺激素受体
 D. Thyroid hormone receptors

答案: B;C
 Answer: B; C

已知抛物线 $\Gamma: y^2 = 4x$ 的焦点为 F, 过 F 作斜率为 k 的直线交 Γ 于 A、B 两点。 (1) 若 $|AB|=10$, 求线段 AB 中点的横坐标; (2) 证明: 在抛物线 Γ 上所有的点中, 顶点到焦点 F 的距离最小;
 It is known that the focus of the parabolic $\Gamma: y^2 = 4x$ is F, and the straight line with a slope of k crosses Γ at points A and B. (1) If $|AB|=10$, find the abscissa of the midpoint of the line segment AB. (2) Prove: Among all the points on the parabolic Γ , the distance from the vertex to the focal point F is the smallest.

答案: (1) AB 中点横坐标为 4; (2) 顶点到焦点距离最小;
 Answer: (1) The abscissa of the midpoint of AB is 4. (2) The distance from the vertex to the focal point is the smallest.

(a) Multiple-choice question within the knowledge dimension

(b) Short-answer question within the knowledge dimension

Figure 4: Example of (a) a multiple-choice question in the knowledge from Biology. (b) A short-answer question in the knowledge from Math. English translations are shown for better readability.

Table 3: Zero-shot average accuracy (%) across six categories in the **knowledge**. The highest accuracy is **bold**, and the second highest is underlined. More results are provided in the Appendix.

Model	Parameters	Access	Creator	FD	HH	SSEM	LP	MH	IIS	Average
Qwen3	8B	Weights	Alibaba	53.02	38.53	36.58	30.17	36.71	37.75	43.86
Qwen3	14B	Weights	Alibaba	36.32	36.78	35.12	27.29	36.82	35.67	35.62
MuduoLLM	14B	Weights	BNU & TAL	28.20	40.82	32.99	36.15	39.11	31.40	33.68
QwQ	32B	Weights	Alibaba	<u>61.25</u>	48.51	42.24	<u>49.90</u>	55.01	47.26	<u>53.87</u>
Seed-OSS	36B	Weights	ByteDance	48.81	<u>50.14</u>	<u>45.34</u>	48.66	61.00	<u>49.56</u>	49.53
Qwen2.5	72B	Weights	Alibaba	19.53	30.95	20.57	13.26	23.86	20.90	22.76
Qwen3	235B (22B active)	Weights	Alibaba	34.24	47.01	36.21	44.26	58.71	46.61	40.82
DeepSeek-V3.1	671B (37B active)	Weights	DeepSeek	31.65	40.65	35.00	29.42	50.54	45.19	36.05
GPT-4o	Undisclosed	API	OpenAI	21.15	26.94	23.92	22.13	34.75	27.13	24.17
Claude-4 Sonnet	Undisclosed	API	Anthropic	41.49	44.29	35.36	27.56	34.86	42.34	40.35
Gemini-2.5 Pro	Undisclosed	API	Google	73.83	<u>55.13</u>	<u>46.68</u>	<u>55.40</u>	<u>60.68</u>	54.16	62.76

instruction-following models (Qwen2.5 (Yang et al., 2024), Qwen3 (Yang et al., 2025)); (b) general reasoning models (QwQ (Team, 2025c), Seed-OSS (Team, 2025b), DeepSeek-V3.1 (Liu et al., 2024b)); and (c) education-specific models (MuduoLLM (from BNU & TAL, 2025)).

Implementation details. In our experimental setup, we evaluate all large language models under both zero-shot and few-shot settings, with few-shot examples (0-, 1-, 3-, and 5-shot) drawn from a separately partitioned development set, distinct from the evaluation set. All open-source models are run using their official code, while closed-source models are accessed via official APIs. We consistently use Gemini-2.5 Pro as the LLM-assisted scoring model, unless otherwise specified.

3.2 MAIN RESULTS

We evaluated all baseline models on OmniEduBench, reporting both per-task category and overall accuracy, as shown in Tables 3 and 4). Results show that in the knowledge dimension, Gemini-2.5 Pro achieves the highest accuracy at 62.78%, while in the cultivation dimension, the reasoning-enhanced version of QWQ performs best with an accuracy of 70.27%. This performance highlights the challenging nature and strong discriminative power of the constructed OmniEduBench.

In the knowledge dimension, it is evident that, except for Gemini-2.5 Pro, closed-source models generally perform worse than open-source models on our OmniEduBench. For example, GPT-4o achieves an accuracy of 24.17%, far below that of Qwen3-8B. This may indicate that the GPT series has relatively weak robustness when handling Chinese education exam-style questions. Meanwhile, model architecture has a significant impact on performance, such as Seed-OOS outperforms the Qwen family by more than 10%. In the cultivation dimension, models generally perform better than in the knowledge dimension, which may be due to the fact that the cultivation tasks mainly consist of multiple-choice questions, making them simpler compared to knowledge tasks with 11 common exam question types. However, differences in performance between different model architectures still exist. Overall, GPT-40 performs the worst in both dimensions, with accuracy largely concentrated around 59.57%, possibly because it has not been specifically optimized for this dimension.

Table 4: Zero-shot average accuracy (%) across six categories in the **cultivation**. The highest accuracy is **bold**, and the second highest is underlined. More results are provided in the Appendix.

Model	Parameters	Access	Creator	TCS	EMH	SIS	CV	PD	TFS	Average
Qwen3	8B	Weights	Alibaba	70.95	66.67	69.16	62.25	70.13	77.20	68.62
Qwen3	14B	Weights	Alibaba	67.79	60.77	63.72	56.20	64.31	71.50	63.60
MuduoLLM	14B	Weights	BNU & TAL	64.42	60.77	63.45	66.14	67.51	64.77	63.96
QwQ	32B	Weights	Alibaba	73.16	<u>68.36</u>	69.84	65.13	71.77	<u>72.02</u>	70.27
Seed-OSS	36B	Weights	ByteDance	70.74	<u>65.30</u>	66.03	62.82	67.12	<u>70.47</u>	67.18
Qwen2.5	72B	Weights	Alibaba	67.89	64.38	65.62	59.51	65.57	67.88	65.34
Qwen3	235B (22B active)	Weights	Alibaba	67.84	61.10	64.54	55.76	64.40	70.47	63.74
DeepSeek-V3.1	671B (37B active)	Weights	DeepSeek	71.58	65.41	69.02	61.96	<u>71.00</u>	77.20	68.55
GPT-4o	Undisclosed	API	OpenAI	61.63	59.57	59.24	55.33	57.71	65.80	59.57
Claude-4-sonnet	Undisclosed	API	Anthropic	71.95	70.05	<u>70.92</u>	64.55	69.25	71.50	70.03
Gemini-2.5-pro	Undisclosed	API	Google	72.26	66.07	<u>70.79</u>	<u>65.71</u>	70.32	67.36	69.14

Table 5: Average accuracy (%) across six categories in one-shot, three-shot, and five-shot settings for the knowledge dimension. The highest accuracy is **bold**, and the second highest is underlined.

Model	Parameters	Access	Creator	FD	HH	SSEM	LP	MH	IIS	Average
<i>One-shot setting</i>										
Qwen3	8B	Weights	Alibaba	52.80	46.45	41.90	29.76	40.20	40.59	41.95
MuduoLLM	14B	Weights	BNU & TAL	27.36	<u>47.79</u>	36.72	34.98	40.74	34.35	36.99
Qwen2.5	72B	Weights	Alibaba	21.42	40.43	27.04	20.96	28.10	22.65	26.77
Qwen3	235B (22B activate)	Weights	Alibaba	<u>37.72</u>	60.79	<u>44.03</u>	45.77	59.59	54.05	50.12
DeepSeek-V3.1	671B (37B activate)	Weights	DeepSeek	30.00	41.73	34.65	30.72	<u>49.67</u>	42.12	38.15
<i>Three-shot setting</i>										
Qwen3	8B	Weights	Alibaba	52.98	46.00	39.74	30.65	39.32	40.85	41.59
MuduoLLM	14B	Weights	BNU & TAL	27.32	<u>46.86</u>	35.79	33.81	39.54	32.42	35.96
Qwen2.5	72B	Weights	Alibaba	21.43	40.86	27.27	20.41	27.12	23.88	26.83
Qwen3	235B (22B activate)	Weights	Alibaba	<u>37.52</u>	60.70	<u>43.40</u>	45.77	59.48	52.57	49.54
DeepSeek-V3.1	671B (37B activate)	Weights	DeepSeek	29.09	41.42	34.02	28.59	48.80	<u>42.06</u>	37.33
<i>Five-shot setting</i>										
Qwen3	8B	Weights	Alibaba	56.86	46.70	39.44	30.65	38.24	42.23	42.35
MuduoLLM	14B	Weights	BNU & TAL	26.93	<u>46.57</u>	36.28	35.74	39.11	34.57	36.53
Qwen2.5	72B	Weights	Alibaba	21.46	41.19	26.96	20.82	26.03	28.56	27.50
Qwen3	235B (22B activate)	Weights	Alibaba	<u>37.40</u>	60.41	<u>44.13</u>	45.77	58.61	55.58	50.32
DeepSeek-V3.1	671B (37B activate)	Weights	DeepSeek	29.39	41.17	32.87	28.45	<u>47.49</u>	38.95	36.39

3.3 ANALYSIS AND FINDINGS

In this section, we further conduct extensive experiments at multiple levels, including few-shot examples, OmniEduBench HARD, and various LLM-assisted scoring methods.

Results in few-shot examples. In Table 5, we present in-context experimental results using different numbers of shots. As the number of shots increases, model performance generally improves; however, the overall gain is limited when considering the average results. We speculate that the drop in accuracy for some models is due to the fact that they have not (or not appropriately) incorporated few-shot examples during the instruction tuning stage. These findings suggest that while few-shot prompting can be beneficial for certain models, its effectiveness strongly depends on the model’s pretraining and instruction tuning strategies. Moreover, the limited average improvement indicates that simply increasing the number of shots may not always lead to substantial gains, highlighting the need for more sophisticated methods to integrate few-shot examples effectively.

Results on OmniEduBench HARD. In Figures 5 and 6, we present the average accuracy of each model on OmniEduBench HARD. OmniEduBench HARD is a subset of OmniEduBench, consisting of the bottom 26% of samples based on model performance, including approximately 1.552K cultivation samples and 7.620K knowledge samples, for a total of 9.172K examples. The experimental results show that: (1) all 11 LLMs exhibit a significant performance drop on OmniEduBench HARD, with even the best-performing model, Gemini, achieving less than 50% accuracy; (2) Qwen2.5-72 performs the worst, significantly lower than the other models, indicating limited capability in handling difficult samples. These findings indicate that further research is needed to enhance LLMs’ ability to generalize and maintain high performance on hard subsets of educational benchmarks.

Results using different LLM-assisted scoring methods. In Table 6, we present the experimental results using different LLM-assisted scoring methods. The performance of the scoring model

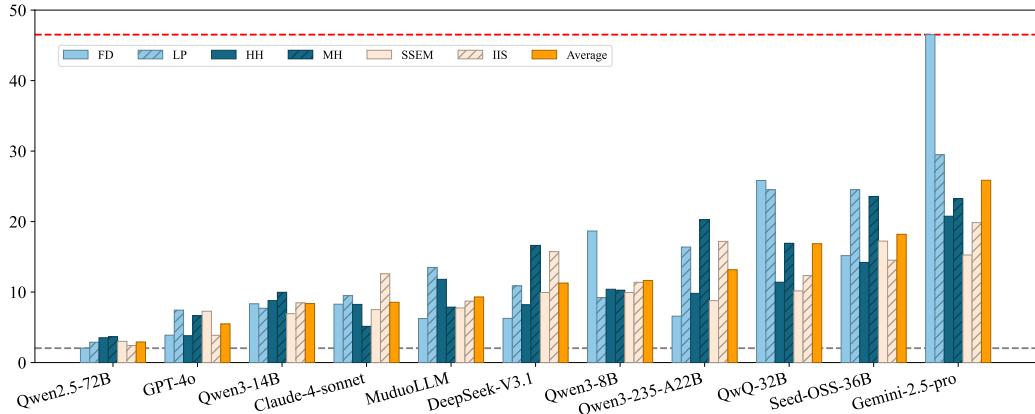


Figure 5: Zero-shot average accuracy (%) on the knowledge dimension of OmniEduBench HARD.

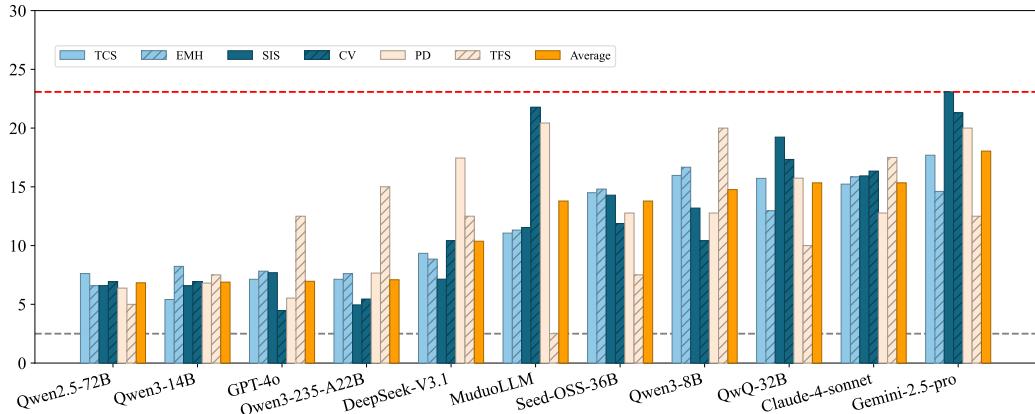


Figure 6: Zero-shot average accuracy (%) on the cultivation dimension of OmniEduBench HARD.

directly affects the evaluation outcomes: higher-quality scoring models provide more accurate assessments, leading to more precise measurements of the evaluated models' capabilities. In this study, we employed three scoring models of varying quality. Overall, GPT-4o performed relatively poorly as a scoring model, failing to accurately evaluate the responses of LLMs. Consequently, the overall effectiveness of LLM-assisted evaluation is reduced when GPT-4o is used, highlighting the critical importance of selecting high-quality scoring models to ensure accurate and meaningful assessments. These findings suggest that the choice of scoring model can substantially influence the perceived performance of evaluated LLMs, and careful selection of scoring models is necessary.

4 RELATED WORK

In this section, we present a comprehensive survey of large language models (LLMs) and benchmarks related to our constructed OmniEduBench, encompassing both English and Chinese datasets.

4.1 LARGE LANGUAGE MODELS

Recently large language models have advanced at an unprecedented pace. Leveraging increasingly sophisticated architectures and ever-larger pretraining corpora, they have continuously pushed the boundaries of performance in language understanding, reasoning, and generation tasks. Researchers have explored various approaches to enhance LLMs' capabilities. For example, Chain-of-Thought prompting (Wei et al., 2022; Team, 2025c;b; Guo et al., 2025; Liu et al., 2024a) has been shown to be highly effective in guiding models to perform step-by-step reasoning for complex problem-solving. In addition, instruction tuning (Dong et al., 2025; Hu et al., 2025; Yang et al., 2024; 2025) and reinforcement learning from human feedback (RLHF) (Ouyang et al., 2022; Schulman et al., 2017) have been widely adopted to align model outputs with human intentions and preferences, enabling

Table 6: Zero-shot average accuracy (%) across six categories in the knowledge using different LLM-assisted scoring methods. The highest accuracy is **bold**, and the second highest is underlined.

Model	Parameters	Access	Creator	FD	HH	SSEM	LP	MH	IIS	Average
<i>Qwen3-A235B-assisted scoring method (Yang et al., 2025)</i>										
QwQ	32B	Weights	Alibaba	<u>61.26</u>	55.82	45.22	<u>50.10</u>	55.66	48.36	<u>56.39</u>
Seed-OSS	36B	Weights	ByteDance	<u>51.16</u>	<u>63.04</u>	<u>51.55</u>	50.38	62.31	<u>57.33</u>	55.49
GPT-4o	Undisclosed	API	OpenAI	23.59	36.77	28.79	23.92	35.73	31.40	28.96
Claude-4 Sonnet	Undisclosed	API	Anthropic	43.54	55.52	40.47	27.77	36.17	47.48	45.34
Gemini-2.5 Pro	Undisclosed	API	Google	75.01	65.67	52.95	56.29	<u>61.44</u>	61.49	67.41
<i>Gemini-2.5 Pro-assisted scoring method (Comanici et al., 2025)</i>										
QwQ	32B	Weights	Alibaba	<u>61.25</u>	48.51	42.24	49.90	55.01	47.26	<u>53.87</u>
Seed-OSS	36B	Weights	ByteDance	48.81	<u>50.14</u>	45.34	48.66	61.00	<u>49.56</u>	49.53
GPT-4o	Undisclosed	API	OpenAI	21.15	26.94	23.92	22.13	34.75	27.13	24.17
Claude-4 Sonnet	Undisclosed	API	Anthropic	41.49	44.29	35.36	27.56	34.86	42.34	40.35
Gemini-2.5 Pro	Undisclosed	API	Google	73.83	55.13	46.68	55.40	<u>60.68</u>	54.16	62.76
<i>GPT-4o-assisted scoring method (Hurst et al., 2024)</i>										
QwQ	32B	Weights	Alibaba	<u>56.61</u>	42.87	37.86	49.28	54.58	37.97	<u>49.26</u>
Seed-OSS	36B	Weights	ByteDance	<u>45.07</u>	<u>47.89</u>	41.94	48.52	60.78	43.54	<u>46.61</u>
GPT-4o	Undisclosed	API	OpenAI	20.38	23.78	22.03	22.06	34.31	23.30	22.51
Claude-4 Sonnet	Undisclosed	API	Anthropic	40.43	41.70	31.89	27.90	35.08	34.79	38.48
Gemini-2.5 Pro	Undisclosed	API	Google	70.15	51.38	44.13	55.88	<u>60.57</u>	46.83	59.49

LLMs to generate responses that are more natural and reliable in open-ended dialogue and creative tasks. Despite these remarkable advances, however, the question of how to comprehensively and effectively evaluate the true capabilities of LLMs remains a critical and open challenge.

4.2 ENGLISH EDUCATION BENCHMARKS

Researchers have proposed a variety of benchmarks to evaluate the capabilities of LLMs, which can be broadly categorized into three types: (1) task-specific evaluations, such as reading comprehension (SQuAD (Rajpurkar et al., 2016)), machine translation (Bojar et al., 2014)), and summarization (Hermann et al., 2015); (2) general knowledge and advanced ability evaluations, for example, the Massive Multitask Language Understanding (MMLU) benchmark (Hendrycks et al., 2021), which collects questions from real-world exams and textbooks to provide a diverse, multi-domain test that effectively probes the breadth and depth of model knowledge. Similarly, the BIG-bench benchmark (Srivastava et al., 2022) comprises 204 diverse tasks; and (3) specialized ability evaluations. In mathematical reasoning, benchmarks such as GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021) assess models’ ability to solve complex multi-step problems. In code generation, HumanEval (Chen et al., 2021) and MBPP (Austin et al., 2021) have become standard benchmarks for measuring programming proficiency. Additionally, datasets such as MT-bench (Zheng et al., 2023) have been introduced to evaluate performance in multi-turn, open-ended dialogues. Despite the significant contributions of these datasets to advancing LLMs evaluation, most of them remain heavily focused on English, with limited coverage of Chinese scenarios.

4.3 CHINESE EDUCATION BENCHMARKS

A series of comprehensive Chinese benchmarks have been proposed. For example, CLUE (Xu et al., 2020b), as an early work, integrates multiple natural language understanding tasks and has become an important reference for evaluating LLMs. Subsequently, benchmarks such as CMMLU (Li et al., 2023) and C-Eval (Huang et al., 2023) collect multi-disciplinary, multi-task questions from Chinese university exams, professional qualification tests, and textbooks, effectively assessing models’ general knowledge and their understanding. Beyond general capability evaluation, researchers have also developed Chinese benchmarks targeting specific advanced skills. For example, in mathematical reasoning, CMATH (Wei et al., 2023) tests models’ abilities to solve complex mathematical problems. Meanwhile, EduBench (Xu et al., 2025) constructs synthetic corpora for the education, but its question types are relatively limited, making it difficult to fully capture models’ Chinese potential. To address this critical gap, we propose OmniEduBench — a comprehensive Chinese education benchmark that uniquely combines knowledge and nurturing dimensions, providing a novel, holistic framework for systematically evaluating LLMs’ potential as educational assistants.

5 CONCLUSIONS, DISCUSSIONS AND LIMITATIONS

In this paper, we present OmniEduBench, a comprehensive Chinese educational benchmark designed to address the limitations of existing Chinese educational evaluation benchmarks. By moving beyond simple knowledge retrieval, the benchmark provides a holistic assessment of LLMs' capabilities across two core dimensions: the knowledge and cultivation dimensions. We conducted extensive experiments on 11 mainstream LLMs, revealing significant performance gaps. While some models performed well on the knowledge dimension, their performance on cultivation tasks dropped substantially, with even the best-performing models trailing human-level performance by nearly 30%. These findings indicate that despite recent advancements in LLM technology, current models still lack the deep reasoning and pedagogical skills necessary to function effectively as educational assistants. We believe OmniEduBench will serve as an important tool for guiding future research. Looking ahead, OmniEduBench plans to explore more complex question types in the cultivation dimension and introduce multimodal educational scenarios, further enhancing the benchmark's role in evaluating and guiding the comprehensive capabilities of LLMs and MLLMs.

Ethics statement. Our constructed OmniEduBench educational benchmark is built from publicly available educational resources as well as authorized private resources permitting open-source use, strictly adhering to copyright and licensing requirements. All data have been systematically processed to remove personally identifiable information (PII) and sensitive content, ensuring privacy and security. The dataset is intended solely for research purposes, aiming to advance the development and evaluation of large language models (LLMs) in educational scenarios.

Reproducibility statement. To ensure reproducibility, we provide detailed descriptions of the dataset construction process, annotation criteria, and experimental settings in both the main paper and the Appendix. The proposed OmniEduBench education dataset, together with preprocessing scripts, evaluation metrics, and model prompts, will be publicly released upon acceptance. All experiments were conducted using standard LLM APIs or open-source checkpoints, with model versions, hyperparameters, and evaluation protocols explicitly documented. This ensures that other researchers can faithfully replicate our results and readily extend the benchmark in future studies.

REFERENCES

- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Ge Bai, Jie Liu, Xingyuan Bu, Yancheng He, Jiaheng Liu, Zhanhui Zhou, Zhuoran Lin, Wenbo Su, Tiezheng Ge, Bo Zheng, et al. Mt-bench-101: A fine-grained benchmark for evaluating large language models in multi-turn dialogues. *arXiv preprint arXiv:2402.14762*, 2024.
- Ondřej Bojar, Christian Buck, Christian Federmann, Barry Haddow, Philipp Koehn, Johannes Leveling, Christof Monz, Pavel Pecina, Matt Post, Hervé Saint-Amand, Radu Soricut, Lucia Specia, and Aleš Tamchyna. Findings of the 2014 workshop on statistical machine translation. In *Proceedings of the Ninth Workshop on Statistical Machine Translation*, pp. 12–58, 2014.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv:1803.05457v1*, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Gheorghe Comanici, Eric Bieber, Mike Schaeckermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blstein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the

- frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.
- Xinpeng Dong, Min Zhang, Didi Zhu, Ye Jun Jian, Zhang Keli, Aimin Zhou, Fei Wu, and Kun Kuang. Ericit: Enhancing robustness by identifying concept tokens in zero-shot vision language models. In *Forty-second International Conference on Machine Learning*, 2025.
- MuduLLM Contributors from BNU and TAL. Muduollm: A high-performance llm for intelligent education solutions. <https://huggingface.co/ERC-ITEA/MuduLLM>, 2025.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638, 2025.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=d7KBjmI3GmQ>.
- Karl Moritz Hermann, Tomas Kociský, Edward Grefenstette, Lasse Espeholt, Will Kay, Mustafa Suleyman, and Phil Blunsom. Teaching machines to read and comprehend. *Advances in neural information processing systems*, 28, 2015.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. Training compute-optimal large language models. *arXiv preprint arXiv:2203.15556*, 2022.
- Hanglei Hu, Yingying Guo, Zhikang Chen, Sen Cui, Fei Wu, Kun Kuang, Min Zhang, and Bo Jiang. Advancing personalized learning with neural collapse for long-tail challenge. In *Forty-second International Conference on Machine Learning*, 2025.
- Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu, Chuancheng Lv, Yikai Zhang, Yao Fu, et al. C-eval: A multi-level multi-discipline chinese evaluation suite for foundation models. *Advances in Neural Information Processing Systems*, 36: 62991–63010, 2023.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Osztrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.
- Haonan Li, Yixuan Zhang, Fajri Koto, Yifei Yang, Hai Zhao, Yeyun Gong, Nan Duan, and Timothy Baldwin. Cmmlu: Measuring massive multitask language understanding in chinese, 2023.
- Aixin Liu, Bei Feng, Bin Wang, Bingxuan Wang, Bo Liu, Chenggang Zhao, Chengqi Dengr, Chong Ruan, Damai Dai, Daya Guo, et al. Deepseek-v2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024a.
- Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024b.
- OpenAI. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35: 27730–27744, 2022.
- Bhrij Patel, Souradip Chakraborty, Wesley A Suttle, Mengdi Wang, Amrit Singh Bedi, and Dinesh Manocha. Aime: Ai system optimization via multiple llm evaluators. *arXiv preprint arXiv:2410.03131*, 2024.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text. *arXiv preprint arXiv:1606.05250*, 2016.

- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*, 2022.
- Liangtai Sun, Yang Han, Zihan Zhao, Da Ma, Zhennan Shen, Baocai Chen, Lu Chen, and Kai Yu. Scieval: A multi-level large language model evaluation benchmark for scientific research. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 19053–19061, 2024.
- Ross Taylor, Marcin Kardas, Guillem Cucurull, Thomas Scialom, Anthony Hartshorn, Elvis Saravia, Andrew Poulton, Viktor Kerkez, and Robert Stojnic. Galactica: A large language model for science. *arXiv preprint arXiv:2211.09085*, 2022.
- Anthropic Team. Claude-4 sonnet, 2025a. URL <https://www.anthropic.com/news/claude-4>. [Online; accessed 23-May-2025].
- ByteDance Seed Team. Seed-oss open-source models. <https://github.com/ByteDance-Seed/seed-oss>, 2025b.
- Qwen Team. Qwq-32b: Embracing the power of reinforcement learning, March 2025c. URL <https://qwenlm.github.io/blog/qwq-32b/>.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *International Conference on Learning Representations, ICLR*, 2019.
- Bin Wang, Chao Xu, Xiaomeng Zhao, Linke Ouyang, Fan Wu, Zhiyuan Zhao, Rui Xu, Kaiwen Liu, Yuan Qu, Fukai Shang, et al. Mineru: An open-source solution for precise document content extraction. *arXiv preprint arXiv:2409.18839*, 2024.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- Tianwen Wei, Jian Luan, Wei Liu, Shuang Dong, and Bin Wang. Cmath: Can your language model pass chinese elementary school math test?, 2023.
- Bin Xu, Yu Bai, Huashan Sun, Yiguan Lin, Siming Liu, Xinyue Liang, Yaolin Li, Yang Gao, and Heyan Huang. Edubench: A comprehensive benchmarking dataset for evaluating large language models in diverse educational scenarios. *arXiv preprint arXiv:2505.16160*, 2025.
- Liang Xu, Hai Hu, Xuanwei Zhang, Lu Li, Chenjie Cao, Yudong Li, Yechen Xu, Kai Sun, Dian Yu, Cong Yu, Yin Tian, Qianqian Dong, Weitang Liu, Bo Shi, Yiming Cui, Junyi Li, Jun Zeng, Rongzhao Wang, Weijian Xie, Yanting Li, Yina Patterson, Zuoyu Tian, Yiwen Zhang, He Zhou, Shaoweihsia Liu, Zhe Zhao, Qipeng Zhao, Cong Yue, Xinrui Zhang, Zhengliang Yang, Kyle Richardson, and Zhenzhong Lan. CLUE: A Chinese language understanding evaluation benchmark. In *Proceedings of the 28th International Conference on Computational Linguistics*, pp. 4762–4772. International Committee on Computational Linguistics, December 2020a.
- Liang Xu, Xuanwei Zhang, Lu Li, Hai Hu, Chenjie Cao, Weitang Liu, Junyi Li, Yudong Li, Kai Sun, Yechen Xu, Yiming Cui, Cong Yu, Qianqian Dong, Yin Tian, Dian Yu, Bo Shi, Jun Zeng, Rongzhao Wang, Weijian Xie, Yanting Li, Yina Patterson, Zuoyu Tian, Yiwen Zhang, He Zhou, Shaoweihsia Liu, Qipeng Zhao, Cong Yue, Xinrui Zhang, Zhengliang Yang, and Zhenzhong Lan. Clue: A chinese language understanding evaluation benchmark, 2020b.

Liang Xu, Anqi Li, Lei Zhu, Hang Xue, Changtai Zhu, Kangkang Zhao, Haonan He, Xuanwei Zhang, Qiyue Kang, and Zhenzhong Lan. Superclue: A comprehensive chinese large language model benchmark. *arXiv preprint arXiv:2307.15020*, 2023.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36:46595–46623, 2023.

Wanjun Zhong, Ruixiang Cui, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang, Amin Saied, Weizhu Chen, and Nan Duan. Agieval: A human-centric benchmark for evaluating foundation models. *arXiv preprint arXiv:2304.06364*, 2023.

A SUPPLEMENTARY MATERIAL

Use of LLMs. In this paper, LLMs were utilized in two primary ways: (1) as auxiliary tools for data cleaning and preliminary quality checks under human supervision. (2) As evaluation targets in benchmark experiments. To ensure data quality, no content directly generated by LLMs was included in the released dataset. During manuscript preparation, LLMs were employed for minor language polishing. All ideas, methodologies, and conclusions are original contributions of authors.

A.1 SUPPLEMENTARY STATISTICS OF OMNIEDUBENCH

In Table 7, we present the bilingual names and abbreviations of all subjects in the knowledge dimension. In Table 8, we present the bilingual names and abbreviations of all subjects in the cultivation dimension. In Tables 9, 10, and 11, we present the detailed data distribution for all 61 subjects.

In Figures 8 and 7, more examples of various questions in the knowledge and cultivation dimensions.

Table 7: Bilingual names and abbreviations of all subject in the knowledge dimension.

Abbreviation	English Name	Chinese Name
MATH	Mathematics	数学
CHEM	Chemistry	化学
BIO	Biology	生物
PHY	Physics	物理
NSCI	Nature & Science	自然与科学
PSTAT	Probability & Statistics	概率论与数理统计
PPHY	Plant Physiology	植物生理学
CS	Computer Science	计算机
BCHEM	Biochemistry	生物化学
OS	Operating Systems	操作系统
AMATH	Advanced Mathematics	高等数学
CNET	Computer Networks	计算机网络
LANG	Chinese Language	语文
GEO	Geography	地理
HIST	History	历史
IART	Introduction to Arts	艺术概论
ILING	Introduction to Linguistics	语言学概论
HSTUD	History Studies / Historiography	历史学
HFA	History of Foreign Art	外国美术简史
IARCH	Introduction to Archaeology	考古学概论
HACL	History of Ancient Chinese Literature	中国古代文学史
HWP	History of Western Philosophy	西方哲学史
POL	Politics	政治
IMOR	Ideology & Morality	思想品德
MGMT	Management	管理学
HRM	Human Resource Management	人力资源管理
TAX	Taxation	税法学
PSCI	Political Science	政治学
MARX	Marxist Theory	马克思主义理论
ELOG	Economic Logic	经济学逻辑能力
NJE	National Judicial Exam	法考真题
CLAW	Criminal Law	刑法学
CVLAW	Civil Law	民法学
LAW	Law / Jurisprudence	法学
TCM	Traditional Chinese Medicine	中医综合
WMED	Western Medicine	西医综合
NURS	Nursing	护理学
IT	Information Technology	信息技术
CEM	Civil Engineering Materials	土木工程材料
EDU	Education	教育学
PSY	Psychology	心理学

Table 8: Bilingual names and abbreviations of all subject in the cultivation dimension

Abbreviation	English Meaning	Chinese Name
<i>Major Categories</i>		
TCS	Thinking & Cognitive Skills	思维与认知能力
EMH	Emotional & Mental Health	情感与心理健康
SIS	Social & Interpersonal Skills	社会与人际交往
CV	Character & Values	品格与价值观
PD	Personalized Development	个性化发展
TFS	Teaching Feedback & Support	教学反馈与支持
<i>Subcategories</i>		
IC	Innovation & Creativity	创新与创造力
PSS	Problem-Solving Skills	问题解决能力
CT	Critical Thinking	批判性思维
GRL	Guided Reflective Learning	反思性学习引导
MA	Metacognitive Abilities	元认知能力
GKT	Guiding Knowledge Transfer	引导知识迁移能力
ER	Emotional Regulation	情绪调控能力
EC	Empathy & Compassion	同理心与共情
SCSE	Self-Confidence & Self-Efficacy	自信心与自我效能感
PR	Psychological Resilience	心理韧性与抗挫力
GM	Growth Mindset	成长型思维
TC	Teamwork & Collaboration	团队协作能力
ECOM	Effective Communication	有效沟通能力
SR	Social Responsibility	社会责任感
RA	Responsibility & Accountability	责任感与担当
IH	Integrity & Honesty	正直与诚信
PLP	Personalized Learning Paths	个性化学习路径
IDL	Interest-Driven Learning	兴趣驱动学习
HT	Heuristic Teaching	启发式教学
CTF	Constructive & Timely Feedback	反馈的建设性与及时性

Table 9: Statistics of OmniEduBench for K-12 in the knowledge dimension.

English Name	Chinese Name	选择题 Multiple choice	多选题 Multiple answer	填空题 Fill-in-the-blank	解答题 Short -answer	复合题 Composite questions	总计 Total
Chinese	语文	350	8	1697	1261	51	3367
Mathematics	数学	527	12	1865	1181	142	3727
Chemistry	化学	274	76	799	477	14	1640
History	历史	67	24	63	211	5	370
Geography	地理	78	31	277	173	4	563
Moral Education	思想品德	14	30	34	56	4	138
Politics	政治	260	241	281	64	12	858
Physics	物理	82	15	178	46	16	337
Biology	生物	115	94	360	124	0	693
Nature Science	自然与科学	8	0	23	22	0	53
Information Technology	信息技术	18	2	14	1	1	36
Total	总计	1793	533	5591	3616	249	11.782K

Table 10: Statistics of OmniEduBench for high, college, and professional schools in the knowledge dimension

English Name	Chinese Name	单选题 Single Choice	多选题 Multiple choice	名词解释 Term explanation	简答题 Short answer	论述题 Essay	案例分析题 Case analysis	填空题 Fill-in- blank	计算题 Calculation	判断题 True/ False	逻辑推理 Logical reasoning	总计 Total
Traditional Chinese Medicine	中医综合	230	317	0	0	0	0	0	0	0	0	547
Chinese Ancient Literary History	中国古代文学史	31	14	20	27	19	0	0	0	0	0	111
Human Resource Management	中国人力资源管理	36	35	9	38	19	6	4	0	0	0	147
Jurisprudence	法学	109	153	0	0	0	0	0	0	0	0	262
Criminal Law	刑法学	158	171	2	3	1	0	0	0	0	0	335
History	历史学	26	0	68	6	17	12	0	0	0	0	129
Civil Engineering Materials	土木工程材料	36	2	67	21	57	2	107	39	3	0	334
A Brief History of Foreign Art	外国美术简史	0	0	61	37	15	14	0	0	0	0	127
Psychology	心理学	146	45	0	35	0	1	0	0	0	0	227
Nursing	护理学	61	41	0	5	0	16	0	0	0	0	123
Operating Systems	操作系统	98	0	0	0	0	59	0	0	0	0	157
Politics	政治	7	47	0	0	0	36	0	0	0	0	90
Pedagogy	教育学	179	0	0	45	39	0	0	0	14	0	277
Advanced Mathematics	高等数学	43	0	0	0	3	0	40	55	0	0	141
Plant Physiology	植物生理学	66	0	0	134	33	48	0	0	0	0	281
Probability and Mathematical Statistics	概率论与数理统计	0	0	0	2	0	74	0	257	0	0	333
Civil Law	民法	91	187	0	0	0	0	0	0	0	0	278
Judicial Practice	司法实践	193	430	0	0	0	0	0	0	0	0	623
Biochemistry	高等生物化学	62	0	0	70	23	0	0	0	0	0	155
Taxation	税法学	19	2	7	30	0	2	0	23	11	0	94
Management	管理学	0	0	0	0	0	0	0	0	0	0	202
Economics	经济学	2	0	0	0	0	0	0	0	0	58	60
Archaeology	考古学概论	0	0	121	0	0	0	0	0	0	0	121
Introduction to Archaeology	考古学概论	57	43	80	45	26	0	0	0	37	0	288
Western Medicine	西医综合	159	119	0	0	0	0	0	0	0	0	278
Western Philosophy	西方哲学史	51	0	0	0	0	0	0	0	0	0	51
Computer Science	计算机科学	137	0	0	0	0	47	0	0	0	0	184
Computer Networks	计算机网络	91	0	0	0	0	50	0	0	0	0	141
Introduction to Linguistics	语言学概论	53	29	24	22	26	18	0	0	0	0	172
Marxist Theory	马克思主义理论	0	0	20	31	13	7	0	0	0	0	71
Total	总计	2141	1635	479	551	291	392	151	374	65	260	6339

Table 11: Statistics of OmniEduBench for 20 subjects in the cultivation dimension.

English Name	Chinese Name	Count
Emotional Regulation Skills	情绪调控能力	325
Innovation & Creativity	创新与创造力	275
Heuristic Teaching	启发式教学	434
Sense of Responsibility & Accountability	责任感与担当	330
Problem-Solving Skills	问题解决能力	288
Team Collaboration Skills	团队协作能力	291
Empathy & Sympathy	同理心与共情	385
Self-Confidence & Self-Efficacy	自信心与自我效能感	358
Constructiveness & Timeliness of Feedback	反馈的建设性与及时性	196
Integrity & Honesty	正直与诚信	371
Psychological Resilience & Anti-Frustration Ability	心理韧性与抗挫力	393
Personalized Learning Paths	个性化学习路径	292
Reflective Learning	反思性学习引导	224
Guiding Knowledge Transfer	知识迁移能力	384
Metacognitive Skills	元认知能力	338
Interest-Driven Learning	兴趣驱动学习	320
Critical Thinking	批判性思维	428
Growth Mindset	成长型思维	393
Social Responsibility	社会责任感	317
Effective Communication Skills	有效沟通能力	139
Total	总计	6.481K

我班上有学生因家庭暴力表现出攻击性行为，其他老师建议开除他。我认为他需要帮助而非惩罚。如何说服同事？请选择最能体现‘同理心与共情’和‘问题解决能力’的回答：
A. 这个孩子本质不坏，他只是用攻击来保护自己。我们应该分析他的行为背后的原因，制定支持计划。
B. 开除他只会让他更恨社会，不如我们联合心理老师做一次行为评估，再决定干预方案。
C. 反正他又不是我亲生的，你们决定吧，别到时候出事怪我没提醒。
D. 我觉得他应该去特殊学校，普通班根本管不了这种孩子。

A student in my class exhibited aggressive behavior due to domestic violence, and other teachers recommended expelling him. I think he needs help rather than punishment. How to convince a colleague? Please choose the answer that best reflects 'empathy and sympathy' and 'problem-solving skills': A. This child is not bad by nature, he just uses aggression to protect himself. We should analyze the reasons behind his behavior and develop a support plan.
B. Expelling him will only make him hate society more, so why not do a behavioral assessment with a psychologist and then decide on an intervention plan.
C. Anyway, he is not my biological son, you decide, don't blame me for not reminding me when something happens.
D. I think he should go to a special school, ordinary classes can't control this kind of child at all.

答案：B

Answer: B

Question within the cultivation dimension

我在背诵古诗词时总觉得枯燥，记不住，也不知道有什么用。怎么才能让传统文化学习变得有意义？请选择最能体现“兴趣驱动学习”和“知识迁移能力”的回答：
A. 你得多抄几遍，熟能生巧。
B. 考试要考，必须死记硬背下来。
C. 尝试把这些诗句用在日记或作文里，比如用“长风破浪会有时”鼓励自己面对考试压力，你会发现它有现实力量。
D. 找一首你喜欢的流行歌，看看有没有引用古诗的歌词，比较一下意境。
I always feel bored when reciting ancient poems, I can't remember them, and I don't know what they are used for. How can we make traditional culture learning meaningful? Please choose the answer that best reflects "interest-driven learning" and "knowledge transfer ability":
A. You have to copy it several times, practice makes perfect.
B. To take the exam, you must memorize it.
C. Try to use these verses in your diary or composition, such as 'There will be times when the wind and waves will be long' to encourage yourself to face the pressure of the exam, and you will find that it has practical power.
D. Find a pop song you like, see if there are lyrics quoted from ancient poems, and compare the artistic conception.

答案：C

Answer: C

Question within the cultivation dimension

Figure 7: Examples of different questions in the knowledge dimension and cultivation dimensions.

名词解释：大晟词派
Definition:Dasheng School of Ci Poetry

答案：词是音乐文学，词风常随乐曲的推移而产生变化，北宋后期音乐机构大晟府的建立，促进了当时词风的变化，有以周邦彦为代表的大晟词派的产生。周邦彦作词以协律为主，长调尤尚铺叙，于音律和谐中求词句的浑雅，极沉郁顿挫之致，北宋的慢词发展到周邦彦的清真词，以知音律、备法度和风格醇雅著称，既无柳永的‘词语坐下’之病，又无苏轼的‘多不协律’之讥，成为后世词人取法的榜样，影响极为深远。

Answer: The establishment of the Dasheng Mansion, a music institution in the late Northern Song Dynasty, promoted the change of the style of words at that time, and the emergence of the Dasheng Poetry School represented by Zhou Bangyan. Zhou Bangyan's lyrics are mainly based on harmony, and the long tone is especially laid out, seeking the elegance of the words and sentences in the harmony of the rhythm, extremely melancholic and frustrated, the slow words of the Northern Song Dynasty developed to Zhou Bangyan's halal words, known for knowing the rhythm, preparing the law and the style of mellow elegance, neither Liu Yong's 'words under the dust', nor Su Shi's ridicule of 'more incongruity', which has become a model for later generations of poets to learn from, and the influence is extremely far-reaching.

Noun explanation question within the knowledge dimension

《西厢记》的戏剧冲突线索和情节发展
The dramatic conflict clues and plot development of "The West Chamber"

答案：西厢记》的戏剧冲突有两条线索一是老夫人与张生、红娘之间的矛盾冲突。前者是主线，后者是辅线。两条线索互相制约，交错展开。第二本中孙飞虎围困普救寺是一个大关目，在此之前，普救、张生与老夫人的矛盾冲突是存在的，表现张、一见钟情之后，两人相互爱慕，希望靠近，与老夫人治家谋生、冰释之间矛盾的矛盾。张生的退缩、崔、张二人以为天遂人愿，不料老夫人变卦，使他们的理想破灭泡影。此后，戏剧冲突发生了一系列变化。张生老夫人的冲突表面化，并且，由于对老夫人的极端不满，他们私下以“非法”的形式结合了，“拷红”是戏剧冲突的又一次转折。崔、张私下结合被老夫人发觉，于是拷问红娘。老夫人与红娘之间发生激烈冲突。红娘抓住老夫人的弱点，勇敢斗争，迫使老夫人不得不承认张、张的关系。但矛盾并没有就此解决，老夫人在承认了崔、张的婚事后，立即提出附加条件张必须应举得官，才能成就婚事。这使得戏剧冲突再起波澜。张生的老夫人的条件，崔、张被迫分离。但崔莺莺却明确提出“此一去得官不得官，便成回来”。老夫人坚持相府不招白衣女婿，莺莺则认为：“但得一个并头莲，强似状元及第二人之间的冲突，已经明显表现为两种婚姻观念的冲突。第五本是戏剧冲突的最后解决。尽管郑恒竭力破坏，张生仍然在得官归后与莺莺成婚，使故事得以大团圆结局。

Answer: There are two clues to the dramatic conflict in "The West Chamber": one is the conflict between the old lady and Zheng Heng on one side and Cui Yingying, Zhang Sheng, and the matchmaker on the other, and the other is the conflict between Cui Yingying, Zhang Sheng, and the matchmaker. The former is the main line, and the latter is the auxiliary line. The two clues restrict each other and unfold in a staggered manner. In the second book, Sun Feihu's siege of the Puji Temple is a big pass. Prior to this, the contradictions and conflicts between Yingying, Zhang Sheng and the old lady were latent. After falling in love with Cui and Zhang at first sight, the two fell in love with each other and hoped to get closer, and the contradiction between the old lady's rigorous management of the family and the frosty conduct, that is, the invisible conflict between Yingying's desire for love and the old lady's maintenance of etiquette and the conflict of the gatekeeper. The appearance of Sun Feihu suddenly intensified the development of the plot. The old lady's wish, Zhang Sheng's retreat from the enemy, Cui and Zhang thought that heaven had fulfilled their wishes, but the old lady changed her mind and their ideals came to naught. Since then, a series of theatrical conflicts have taken place, and the conflict between Cui, Zhang and the old lady has surfaced, and, due to extreme dissatisfaction with the old lady, they have been united in private in an "illegal" form, "Torture Red" is another major turn point. The dramatic conflict between the old lady and the matchmaker. The matchmaker leased the old lady's weakness and fought bravely, forcing the old lady to admit their relationship between Cui and Zhang. The matchmaker realized the marriage of Cui and Zhang, and there was a secret conflict between the old lady and the matchmaker. After the old lady agreed to the marriage of Cui and Zhang, she immediately proposed additional conditions: Zhang Sheng must be an official in order to achieve the marriage, which made the dramatic conflict make waves again. Zhang Sheng agreed to the old lady's conditions, and Cui and Zhang were forced to separate. But Cui Yingying clearly proposed that "if you go to an official position this time, you will come back when you are sick". The old lady insisted that the prime minister's mansion did not recruit a son-in-law in white, and Yingying believed: "But there is a parallel lotus, which is stronger than the conflict between the champion and the second person, which has obviously manifested itself as a conflict between the two marriage concepts. The fifth book is the final resolution of the dramatic conflict. Despite Zheng Heng's efforts to destroy it, Zhang Sheng still married Yingying after returning from the official position, so that the story could end happily.

Essay Question within the knowledge dimension

我班上有学生因家庭暴力表现出攻击性行为，其他老师建议开除他。我认为他需要帮助而非惩罚。如何说服同事？请选择最能体现‘同理心与共情’和‘问题解决能力’的回答：

- A.这个孩子本质不坏，他只是用攻击来保护自己。我们应该分析他的行为背后的原因，制定支持计划。
- B.开除他只会让他更恨社会，不如我们联合心理老师做一次行为评估，再决定干预方案。
- C.反正他又不是我亲生的，你们决定吧，别到时候出事怪我没提醒。
- D.我觉得他应该去特殊学校，普通班根本管不了这种孩子。

A student in my class exhibited aggressive behavior due to domestic violence, and other teachers recommended expelling him. I think he needs help rather than punishment. How to convince a colleague? Please choose the answer that best reflects 'empathy and empathy' and 'problem-solving skills'. A. This child is not bad by nature, he just uses aggression to protect himself. We should analyze the reasons behind his behavior and develop a support plan.

- B. Expelling him will only make him hate society more, so why not do a behavioral assessment with a psychologist and then decide on an intervention plan.
- C. Anyway, he is not my biological son, you decide, don't blame me for not reminding me when something happens.
- D. I think he should go to a special school, ordinary classes can't control this kind of child at all.

答案：B

Answer: B

Question within the cultivation dimension

周经理近来很不顺心，各部门都向人力资源部门要人，可一时哪有那么多合适的人啊？这种情况在一年中已出现了三次，周经理不明白是这些部门发了疯，还是自己的工作出了错。为了减轻工作压力，周经理独自来到熟悉的酒楼用餐，无意间听到了一段酒楼经理和顾客的对话。一名顾客一脸不悦地向酒楼经理抱怨：“前天我和家人来用餐，一致认为青椒童子鸡最好吃，当时青椒童子鸡是限量供应特色菜；今天我专门请同事来尝，还特地起了早。不想今天青椒童子鸡成了限时特价菜。害得我被同事嘲笑了一番，说我赶早是为了请大家吃便宜菜。你听，他们还在包间里笑呢。”旁的周经理不禁失笑，一个不知道外面供应什么、一个不知道自己供应什么、不出乱才怪。但转念一想，自己不正也犯着同样的错误吗？一方面不清楚公司内部的人员情况，每次招人都措手不及；一方面也不知道劳动力市场供给情况，常常一时招不到合适的工人。原先的嘲笑变成了自嘲，但这顿晚餐让周经理觉得非常满意。请问：(1)周经理从对话中悟出了什么？(2)你认为周经理回到公司后，会如何解决他所面临的问题？

Manager Zhou has been very unhappy recently, and all departments are asking for people from the human resources department, but how can there be so many suitable people at one time? This situation has occurred three times in a year, and Manager Zhou doesn't understand whether these departments are crazy or his work is wrong. In order to reduce work pressure, Manager Zhou came to a familiar restaurant alone to eat. By chance, he heard a conversation between a restaurant manager and a customer. A customer said: "The green pepper童子 chicken was the most delicious, and the green pepper童子 chicken was a limited supply special dish today; Today I specially invited my colleagues to taste it, and I rushed early. I didn't want the green pepper童子 chicken to become a limited-time special dish today, which caused me to be ridiculed by my colleagues, saying that I rushed to treat everyone to cheap food, so it was strange that there was no chance. But on second thought, am I not making the same mistake? On the one hand, I don't know the personnel situation within the company, and I caught off guard every time I recruit people. On the one hand, it is also unclear about the supply of the labor market, and it is often impossible to recruit suitable workers. The original laughter turned into self-criticism. After this dinner, Manager Zhou felt very satisfied. (1) What did Manager Zhou learn from the conversation? (2) How do you think Manager Zhou will solve the problems he faces after returning to the company?

答案：(1)悟出了人力资源规划的重要性。(5分)\n(2)制定企业的人力资源规划，首先确认现阶段企业经营战略；(2分)其次对现有人力资源进行盘点；(2分)再次对人力资源的需求和供给进行预测，确定人力资源的净需求；(2分)最后执行规划。(2分)并对规划实施监控与评估。(2分)

Answer: (1) Realize the importance of human resource planning. (5 points)\n(2) Formulate the human resource plan of the enterprise, first confirm the current business strategy of the enterprise; (2 points) secondly, take stock of existing human resources; (2 points) Forecast the demand and supply of human resources again, and determine the net demand for human resources; (2 points) Final implementation plan. (2 points) and implement monitoring and evaluation of the plan. (2 points)

Case analysis question within the knowledge dimension

Figure 8: Examples of different questions in the knowledge dimension and cultivation dimensions.